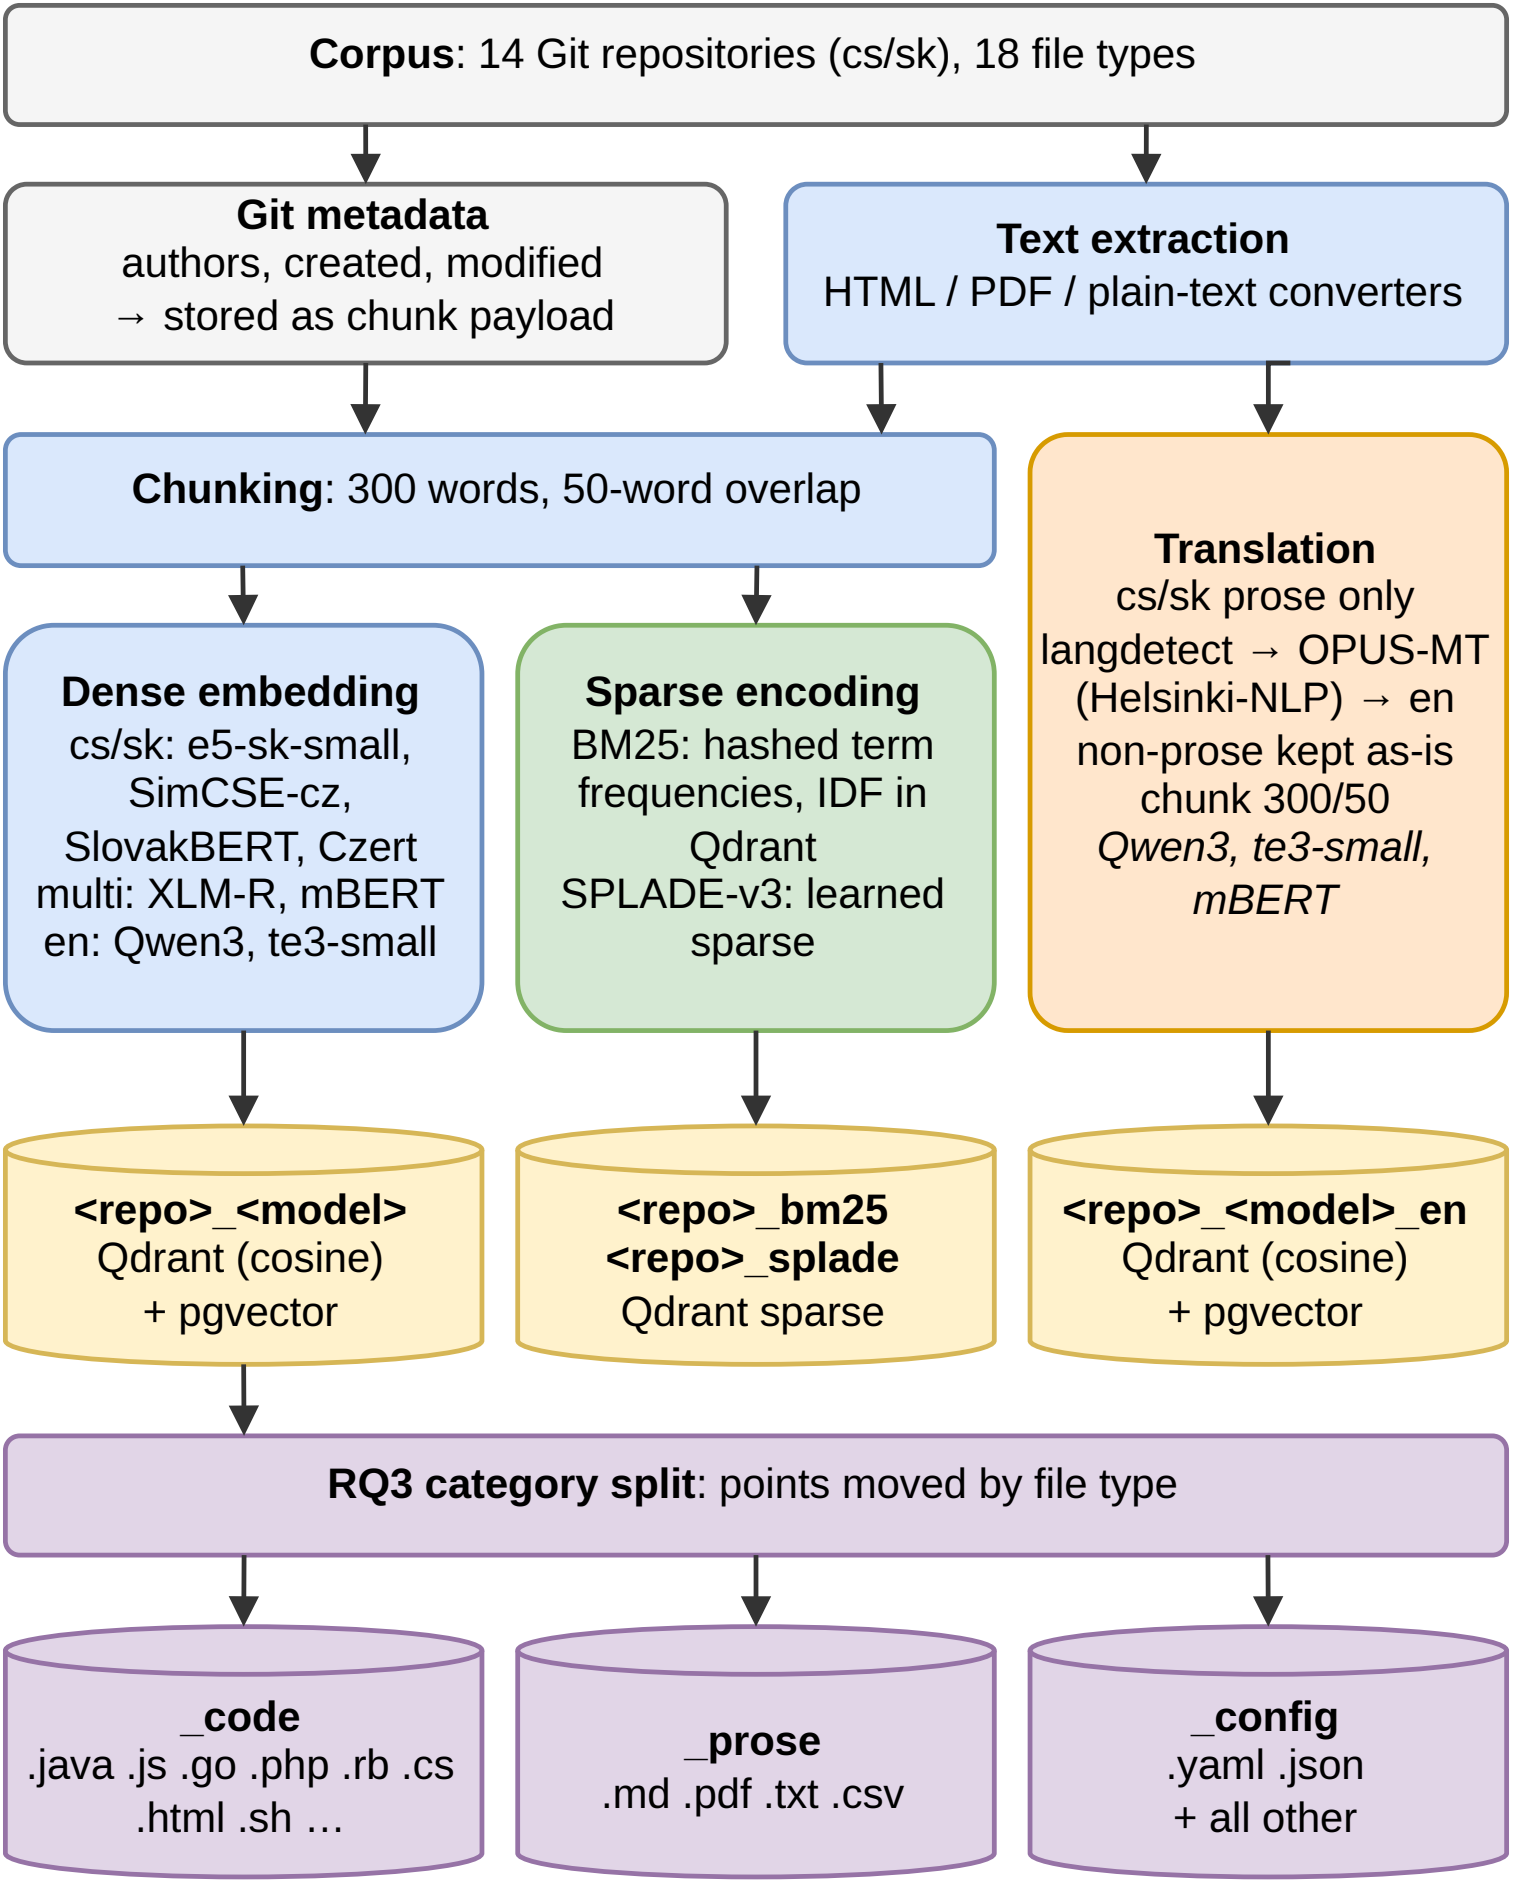
Corpus: 14 Git repositories (cs/sk), 18 file types
Git metadata
authors, created, modified
→ stored as chunk payload
Text extraction
HTML / PDF / plain-text converters
Chunking: 300 words, 50-word overlap
Translation
cs/sk prose only
langdetect → OPUS-MT (Helsinki-NLP) → en
non-prose kept as-is
chunk 300/50
Qwen3, te3-small, mBERT
Dense embedding
cs/sk: e5-sk-small, SimCSE-cz, SlovakBERT, Czert
multi: XLM-R, mBERT
en: Qwen3, te3-small
Sparse encoding
BM25: hashed term frequencies, IDF in Qdrant
SPLADE-v3: learned sparse
<repo>_<model>
Qdrant (cosine)
+ pgvector
<repo>_bm25
<repo>_splade
Qdrant sparse
<repo>_<model>_en
Qdrant (cosine)
+ pgvector
RQ3 category split: points moved by file type
_code
.java .js .go .php .rb .cs .html .sh …
_prose
.md .pdf .txt .csv
_config
.yaml .json
+ all other