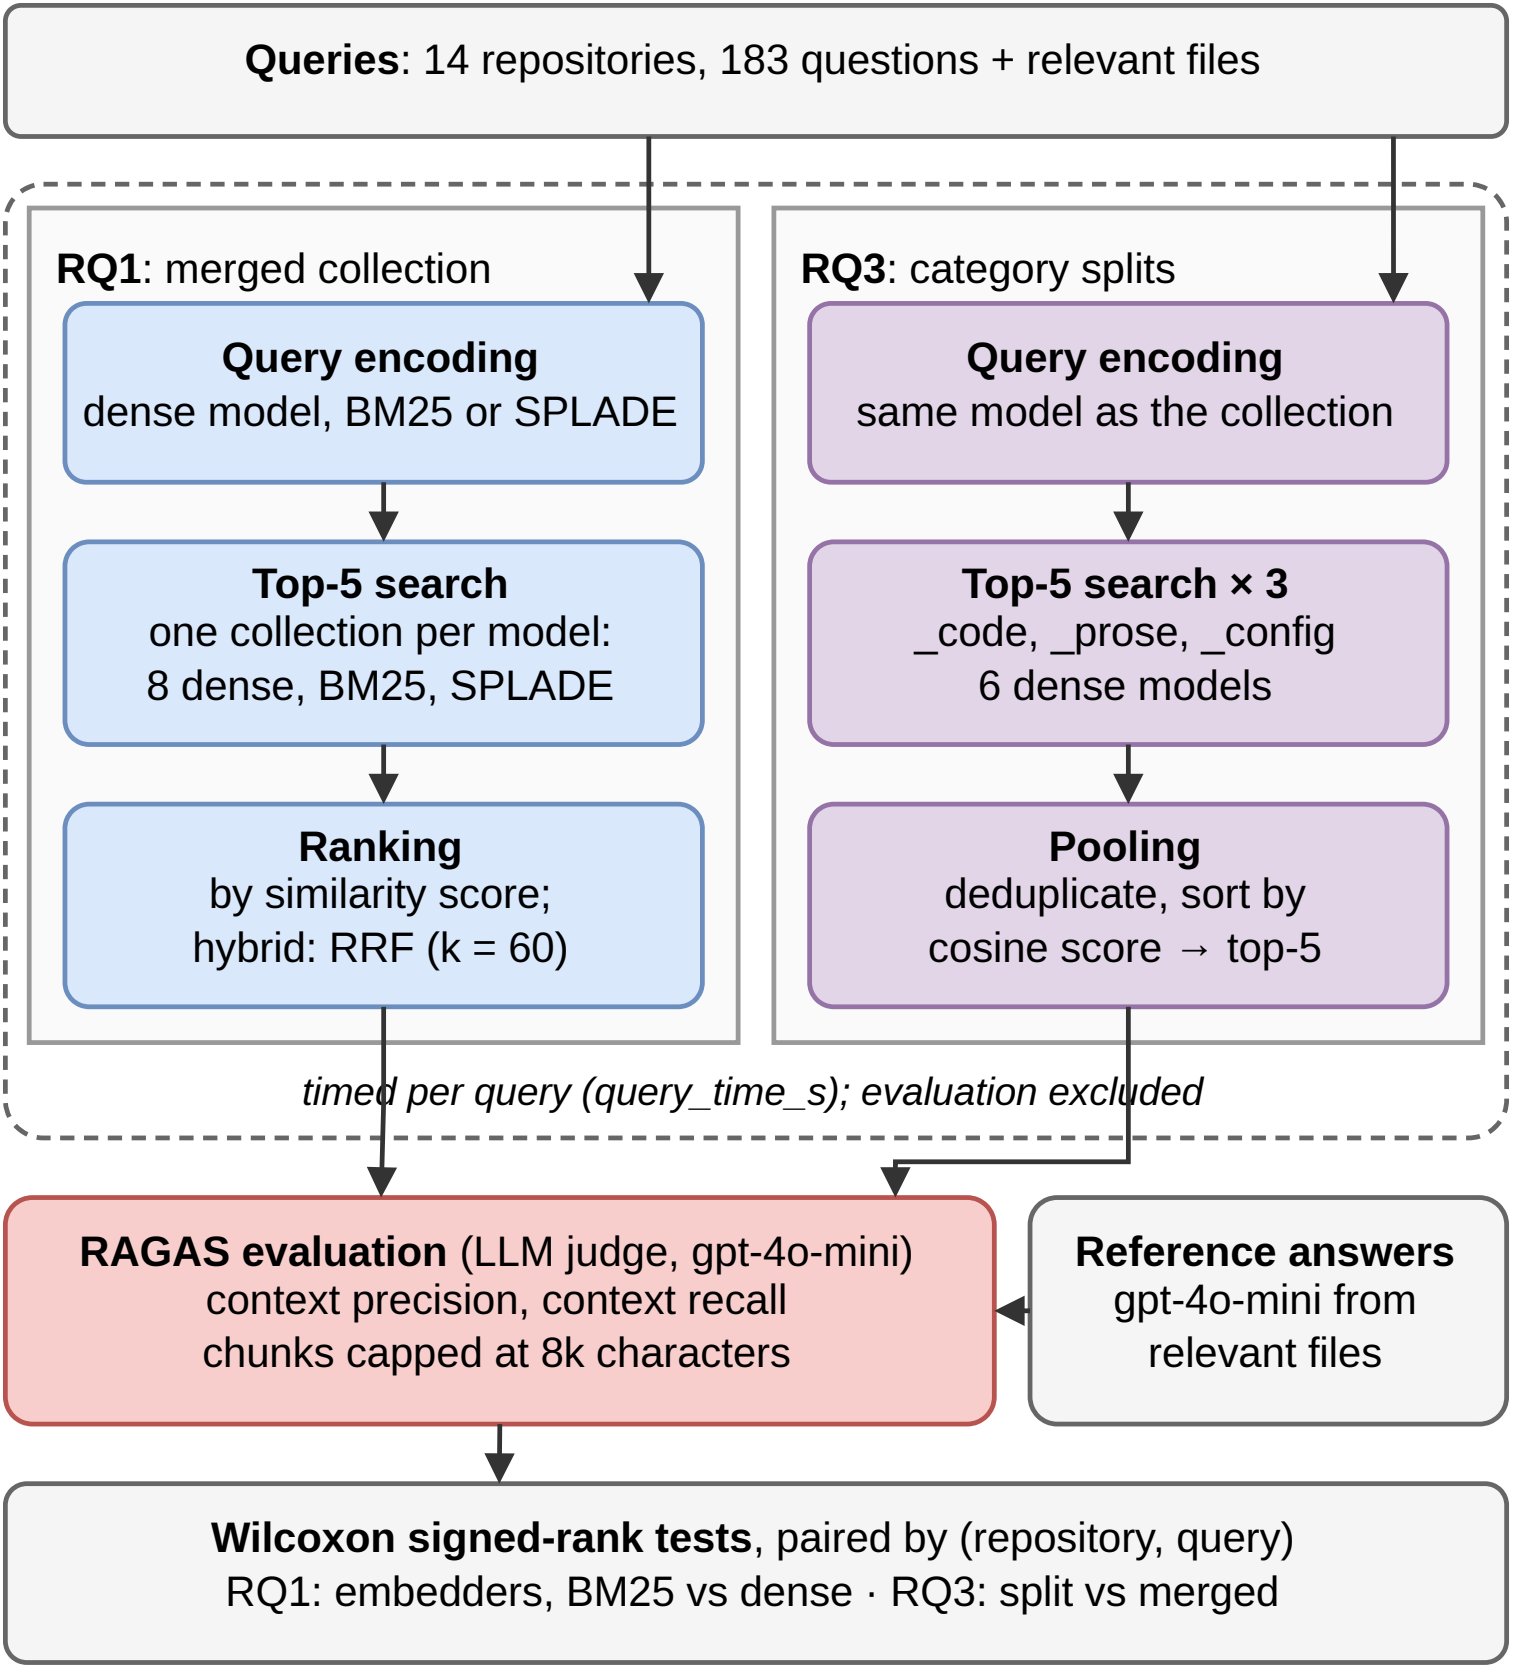
Queries: 14 repositories, 183 questions + relevant files
RQ1: merged collection
Query encoding
dense model, BM25 or SPLADE
Top-5 search
one collection per model:
8 dense, BM25, SPLADE
Ranking
by similarity score;
hybrid: RRF (k = 60)
RQ3: category splits
Query encoding
same model as the collection
Top-5 search × 3
_code, _prose, _config
6 dense models
Pooling
deduplicate, sort by
cosine score → top-5
timed per query (query_time_s); evaluation excluded
RAGAS evaluation (LLM judge, gpt-4o-mini)
context precision, context recall
chunks capped at 8k characters
Reference answers
gpt-4o-mini from
relevant files
Wilcoxon signed-rank tests, paired by (repository, query)
RQ1: embedders, BM25 vs dense · RQ3: split vs merged